

Yanggan Gu

yanggangu@outlook.com | [Google Scholar](#) | [Homepage](#) | [Github](#) |

Education

Soochow University [<i>Project 211</i>], Suzhou, China.	2022/09 - 2025/06
M.S. in Computer Science and Technology.	GPA 3.7 (Top-5%)
Guangdong University of Technology , Guangzhou, China.	2018/09 – 2022/06
B.S. in Software Engineering.	GPA 4.0 (Top-1%)

Research Experience

Model Fusion in LLMs , <i>Research Assistant</i>	2025/04 - Now
Department of Computing, Hong Kong Polytechnic University.	Advisor: Prof. Hongxia Yang
Research Topic: Model Fusion, Alignment of LLMs, Reasoning in LLMs	
Trustworthy LLMs , <i>Research Assistant</i>	2024/06 - 2025/03
AI Thrust, The Hong Kong University of Science and Technology (Guangzhou).	Advisor: Prof. Xuming Hu
Research Topic: Alignment of LLMs, Factual Reasoning in LLMs, Knowledge Distillation	
Syntactic Parsing , <i>Master’s Researcher</i>	2022/09 - 2025/06
School of Computer Science and Technology, Soochow University.	Advisor: Prof. Zhenghua Li
Research Topic: Linguistic Structure Prediction, Text2SQL (maintaining SeSQL dataset)	

Main Publications

InfiFPO: Implicit Model Fusion via Preference Optimization in Large Language Models [Code], <i>First Author</i> .	
Keywords: Alignment, Model Fusion, LLMs.	NeurIPS 2025 Spotlight (CCF-A)
Intro: This paper introduces InfiFPO, a novel preference optimization method for implicit model fusion of LLMs. Unlike prior fusion approaches that mostly rely on supervised fine-tuning or discard probability information during preference alignment, InfiFPO synthesizes multi-source probabilities at the sequence level to replace the reference model in Direct Preference Optimization. Experiments across 11 benchmarks show that InfiFPO consistently surpasses existing model fusion and preference optimization techniques, boosting the Phi-4 model’s average performance from 79.95 to 83.33, with notable improvements in mathematics, coding, and reasoning tasks.	
Capturing Nuanced Preferences: Preference-Aligned Distillation for Small Language Models [Code], <i>First Author</i> .	
Keywords: Alignment, Distillation, Trustworthy LLMs.	ACL 2025 Findings (CCF-A)
Intro: This paper introduces the Preference-Aligned Distillation (PAD) framework to improve the alignment of small language models (SLMs) with human preferences by modeling the full preference distribution rather than using simple pairwise comparisons. Through experiments across multiple benchmarks, the proposed PAD approach is shown to outperform existing methods, offering more precise preference signals for better response generation in SLMs.	
Reasoning Factual Knowledge in Structured Data with Large Language Models [Code], <i>Co-first Author</i> .	
Keywords: Trustworthy LLMs, Benchmark, Structured Facts.	ACL 2025 Findings (CCF-A)
Intro: This paper develops a benchmark, called <i>StructFact</i> , to evaluate the reasoning abilities of large language models (LLMs) over structured factual knowledge, addressing challenges such as heterogeneity, topological interdependencies, and sparsity in structured data. The benchmark consists of 13,407 factual queries across five tasks and examines how LLMs handle fresh, structured data in real-world applications, offering insights into their limitations and potential improvements for knowledge-sensitive tasks.	

Honors and Awards

• 1st place in the CCL-2023 Chinese AMR Parsing Evaluation Task. [Report]	2023/08
• Top 10 Innovative Figures of Guangdong University of Technology. [Sohu News]	2021/12
• National Encouragement Scholarship. [Link]	2021/11

Language

- **English:** TOEFL 91/120 (Reading 27/30 | Listening 22/30 | Speaking 22/30 | Writing 20/30)